
Perceptive Multimodal Generative AI

Rithani M.

Department of

Computer Science and Engineering

Amrita School of Computing

Amrita Vishwa Vidyapeetham

Chennai, India

m_rithani@ch.amrita.edu

Sidesh Sundar S.

Department of

Computer Science and Engineering

Amrita School of Computing

Amrita Vishwa Vidyapeetham

Chennai, India

sidesh260303@gmail.com

Vimal Dharan N.

Department of

Computer Science and Engineering

Amrita School of Computing

Amrita Vishwa Vidyapeetham

Chennai, India

vimalvimal1293@gmail.com

Abstract—Generative AI has led to impressive advancements in text generation, image synthesis, and multimodal interactions. However, many existing systems are confined to single-modality outputs, which limit their effectiveness in creative workflows. In this paper, we introduce Perceptive Multimodal Generative AI (PMG-AI), a powerful and adaptive framework that integrates deep learning, GANs, NLP, OCR, and diffusion models to support end-to-end multimodal content creation, modification, and editing. PMG-AI empowers users to generate diverse outputs across multiple modalities—ranging from text-to-image synthesis, comic creation, and image style transfer to 2D-to-3D conversion and hyper-interactive content generation. The core of the system lies in a CLIP-based encoder that transforms user prompts into latent space representations, which are then refined using latent diffusion models to produce semantically accurate and visually compelling content. Through its dual-path generation mechanism, PMG-AI produces both static and animated content. Its selective inpainting module facilitates localized editing of images, ensuring high-quality results without disrupting the overall structure. Additionally, PMG-AI integrates OCR to simplify text extraction and repurposing from images, opening up new possibilities for content reuse. By enabling 2D-to-3D reconstruction with advanced voxel-based modeling and depth estimation, PMG-AI allows for the creation of 3D assets for AR/VR, digital design, and simulation. This unified, scalable platform caters to creative professionals in fields such as digital art, education, marketing, virtual storytelling, and design, offering innovative tools to amplify the creative process.

Index Terms—Generative AI, Multimodal Content Creation, Text Generation, Image Synthesis, Comic Generation, Image Style Transfer, Optical Character Recognition (OCR), Natural Language Processing (NLP), Model Personalization, AI Content Detection, Deep Learning(DL).

I. INTRODUCTION

Generative AI has significantly transformed creative industries by enabling the automatic generation of text, images, and even voice. While models like GPT-4 for text generation and StyleGAN for image creation have achieved remarkable success in their respective domains, they typically operate in silos, limiting their use in more complex and collaborative workflows. This paper presents Perceptive Multimodal Generative AI (PMG-AI), an integrated platform designed to enhance the capabilities of generative models across multiple modalities,

overcoming the limitations of single-modality systems. PMG-AI is a powerful and adaptive framework that supports a wide array of creative applications such as comic generation, image style transfer, text-to-image synthesis, and even 2D-to-3D content generation. By incorporating a CLIP-based encoder and latent diffusion models, the system generates semantically aligned and visually compelling images based on user prompts. Additionally, PMG-AI features a dynamic generation mechanism that enables both static image creation and animated content, allowing users to explore new possibilities in visual storytelling. A mask-based inpainting module powered by stable diffusion offers selective image editing, while semantic conditioning and cross-attention mechanisms ensure the outputs remain realistic, coherent, and aligned with user intent. OCR integration further enriches the system by facilitating text extraction from images for repurposing, and the system's 2D-to-3D conversion capabilities allow for the transformation of images into high-quality 3D models, expanding its potential for use in virtual reality, simulation, and digital design.

II. LITERATURE REVIEW

Multimodal medical image fusion has gained significant importance in clinical diagnosis, offering improved understanding through the fusion of complementary features of various modalities. Moghtaderi et al. [1] presented an adaptive image decomposition scheme using multilevel guided edge-preserving filtering, which greatly improves contrast and structural details. Das et al. [2] presented a comprehensive review of current methods, databases, and measures of quality, thereby outlining gaps and suggesting future research in the direction of more interpretable and precise fusion methods. Alzahrani [3] presented a new method combining a modified discrete wavelet transform (DWT) with an arithmetic optimization algorithm, which greatly preserves texture and detail between modalities.

Zhang and Wang [4] also suggested a hybrid fusion framework that combines several strategies to improve diagnostic clarity and image consistency. Gu et al. [5] presented AdaFuse, a Fourier-based spatial-frequential attention mechanism that preserves high-frequency image details, thus outperforming baseline models in visual and numerical assessment. Li and

Chen [6] also pushed the boundaries with MedFusionGAN—a GAN-based unsupervised fusion model to fuse CT and MRI scans for better radiotherapy planning. HEALNet, presented by Hemker et al. [7], demonstrated strong multimodal learning by combining heterogeneous biomedical data into a common latent space even in the presence of missing modalities.

Zhou et al. [8] pushed the boundary of edge learning and multiscale feature representation by using a dilated residual attention network to enable the faster and more accurate fusion. Zhou et al.’s [9] Hi-Net introduced a hybrid network for MR image synthesis through modality-specific and shared feature learning. Hill and Hawkes [10] also introduced the key insights into medical image registration and early image computing methods that have influenced modern fusion techniques greatly.

Concurrent with fusion developments, adaptive multimodal image generation has witnessed significant growth with emphasis on integrating various data types—text, style, image, and reference input—for homogeneous visual synthesis. Hu et al. [11] introduced Instruct-Imagen, a multimodal prompt and retrieval-augmented fine-tuning-based image generation system for task-specific output. Li et al. [12] proposed UNIMO-G, a conditional diffusion model accepting both text and image inputs via a Multimodal Large Language Model (MLLM), thereby generating high-quality outputs from sophisticated prompts.

Sun et al. [13] presented Emu, a transformer-based model that can process interleaved image-text-video sequences on a range of tasks such as captioning and visual question answering (VQA). Koh et al. [14] presented an embedding fusion method that combines frozen image and text models to enable multimodal interaction and thereby text-to-image generation and retrieval with ease. Tan et al. [15] developed a generative adversarial network (GAN) model that combines text, image, and style data to generate stylized and coherent visual content.

Adaptive fusion in medical imaging was also studied in the context of a deep pyramidal residual learning network [16], using hierarchical and residual learning methods to enhance fusion quality. Warner et al. [17] conducted a thorough review of multimodal machine learning within clinical systems, such as fusion, co-learning, and translation processes. Further, there was another adaptive model [18] that highlighted the use of trilinear fusion strategies for VQA, stressing the efficiency of collaborative attention methods.

A recent systematic review [19] has described the advancement in multimodal image fusion using deep learning, attention mechanisms, and transformer models, thus pointing out the fusion of multiple input modalities to achieve more stable results. Moreover, a subsequent study [20] re-explored adaptive edge-preserving decomposition methods for fusion, thus pointing out the significance of maintaining semantic and structural information in medical images.

III. METHODOLOGY

In our study, we propose an enhanced version of ChatGPT to produce consistent and insightful visual responses—particularly for color enhancement tasks—we present a pure version of the system in this work, as well as side defense. Artificial intelligence and image processing techniques are used in advanced modeling to produce meaningful conversations and accurate results. It has also improved a lot of performance features, and much more. This method generates model outputs based on perception and context by utilizing artificial adversarial networks (GAN), attention strategies, and cognitive loss functions.

A. TEXT TO IMAGE GENERATION

The framework **Perceptive Multimodal Generative AI (PMG-AI)** translates natural language descriptions into high-quality static or animated images while supporting region-specific image refinement through inpainting. The system begins with text input tokenization via Byte Pair Encoding (BPE) in the CLIP model, followed by encoding through a Transformer to produce a fixed-dimensional semantic vector $E_t \in \mathbb{R}^d$, representing the text’s meaning. The self-attention mechanism in CLIP enables contextual embeddings, guiding downstream generative processes.

The semantic vector E_t conditions a Latent Diffusion Model (LDM) operating in a lower-dimensional latent space \mathcal{Z} , allowing more computational efficiency and better semantic density. The forward diffusion process gradually corrupts a clean latent z_0 by adding Gaussian noise, with the reverse process denoising it back to a clean image using a U-Net neural network. The denoising process is trained with a mean squared error loss:

$$\mathcal{L}_{\text{DDPM}} = \mathbb{E}_{z_0, \epsilon, t} \left[\|\epsilon - \epsilon_\theta(z_t, t, E_t)\|^2 \right]$$

The denoised latent \hat{z}_0 is decoded back into an image using a pretrained VQGAN or autoencoder decoder:

$$\hat{x} = D(\hat{z}_0)$$

For animated content, the system generates multiple images by varying latent initializations or prompt embeddings. Temporal smoothness is ensured by duplicating frames, forming an animation sequence that can be saved in GIF or MP4 format.

PMG-AI also supports **interactive editing** via masked inpainting. A user-defined binary mask M specifies the region to regenerate, and the rest of the image remains intact:

$$x_{\text{final}} = M \odot \hat{x}_{\text{edit}} + (1 - M) \odot x$$

This allows selective alterations, such as object replacement or background changes, without affecting the entire image.

After generation, images are transformed into RGB format for visualization, and optional post-processing (e.g., super-resolution) can be applied to enhance image quality.

1) Token to Latent Diffusion: The *Token to Latent Diffusion* process bridges the text prompt's semantic understanding to visual synthesis. First, a textual prompt T is encoded into a high-dimensional vector E_t using CLIP. The corresponding image is encoded into a latent space z_0 using a pretrained VQGAN or autoencoder. This latent space operates at reduced dimensionality Z , ensuring efficiency without losing critical semantic detail.

$$z_0 = E_{\text{img}}(x), \quad z_0 \in \mathbb{R}^{C \times H' \times W'}$$

In the forward diffusion process, Gaussian noise is added to the clean latent representation over several timesteps. The reverse process is modeled using a U-Net, which learns to denoise the latent and reconstruct the image.

The **reverse diffusion** process iteratively refines the noisy latent representation:

$$z_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(z_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \cdot \epsilon_\theta(z_t, t, E_t) \right) + \sigma_t \cdot \mathcal{N}(0, \mathbf{I})$$

2) Image Generation Process: The *Image Generation Process* in PMG-AI starts with the preparation of a paired dataset of text and images. Each image is resized and normalized, and the caption is tokenized using BPE. The image is then encoded into a latent representation z_0 , and the semantic text embedding E_t is generated from the prompt.

$$E_t = \text{CLIP}_{\text{text}}(T), \quad E_t \in \mathbb{R}^d$$

Once both the text and image are encoded, the Latent Diffusion Model (LDM) is trained to condition on these embeddings, refining the generated image through a diffusion process.

3) Stable Diffusion Training: Forward and Reverse Denoising: In **Stable Diffusion**, the image is corrupted with Gaussian noise over several timesteps. The model learns the reverse process to recover the original latent image. The denoising network ϵ_θ is trained to predict the noise and is conditioned on the semantic text embedding E_t . The reverse denoising is trained using:

$$\mathcal{L}_{\text{DDPM}} = \mathbb{E}_{z_0, \epsilon, t} \left[\|\epsilon - \epsilon_\theta(z_t, t, E_t)\|^2 \right]$$

Once trained, the model can generate clean latents and produce the final image.

4) Sampling with Stable Diffusion: At inference time, the system takes only the text prompt T , which is encoded into E_t . A random latent z_T is sampled, and the model iteratively denoises it back to the original z_0 , which is then decoded to generate the final image.

$$\hat{x} = D(\hat{z}_0)$$

5) Animated Image Generation: For **animated content**, multiple latent vectors are sampled and processed sequentially. Each frame is denoised and decoded into a generated image, and the sequence is smoothed and exported as a GIF or MP4.

6) Masked Inpainting: Region-Specific Regeneration: PMG-AI allows targeted editing using **binary masks** to regenerate selected regions of an image. The masked area is processed separately while retaining the surrounding content. This allows users to refine specific details without disturbing the overall image:

$$x_{\text{final}} = M \odot \hat{x}_{\text{edit}} + (1 - M) \odot x$$

7) Final Output and Enhancement: After generation, the image is optionally enhanced using **super-resolution**, **CLIP-based re-ranking**, and other techniques to improve visual quality. The final image is saved in standard formats like PNG or JPEG, and the process can be seamlessly integrated into creative workflows.

B. 2D to 3D Reconstruction

Converting a 2D image to a 3D model is a complex challenge in computer vision, as it requires inferring hidden spatial and geometric properties. This task is crucial for applications like augmented reality (AR), virtual reality (VR), robotics, and 3D content creation. PMG-AI introduces a modular approach combining deep learning with classical computer vision to accurately predict 3D shapes from 2D images. The framework integrates voxel-based modeling, keypoint estimation, depth estimation, and mesh extraction, enabling both dense and sparse 3D representations.

The pipeline begins with image preprocessing, resizing, and normalizing to prepare the 2D image. The next step involves expanding the image into a pseudo-volumetric tensor, with an added depth dimension to simulate the 3D structure. The resulting tensor $\mathbb{R}^{1 \times 1 \times D \times H \times W}$ is passed through a 3D convolutional encoder-decoder network, producing a voxel occupancy grid $\hat{v} \in [0, 1]^{D \times H \times W}$. The model is trained with binary cross-entropy loss:

$$\mathcal{L}_{\text{voxel}} = - \sum_{i=1}^D \sum_{j=1}^H \sum_{k=1}^W [v_{ijk} \log(\hat{v}_{ijk}) + (1 - v_{ijk}) \log(1 - \hat{v}_{ijk})]$$

To extract the surface mesh, the Marching Cubes algorithm is applied to the voxel grid, generating a 3D mesh saved in formats like .OBJ or .PLY. This mesh is suitable for 3D rendering and simulation.

Additionally, sparse 3D reconstruction is performed through 2D keypoint regression. The system predicts 2D keypoints $K_{2D} \in \mathbb{R}^{N \times 2}$ from the image and assigns synthetic depth values to convert them into 3D keypoints $K_{3D}^{(i)} = [K_{2D}^{(i)}, z_i]$. The keypoint regression is trained using mean squared error loss:

$$\mathcal{L}_{\text{keypoints}} = \frac{1}{N} \sum_{i=1}^N \|K_{2D}^{(i)} - \hat{K}_{2D}^{(i)}\|^2$$

Depth estimation is performed using edge detection (Canny) and Gaussian filtering, refining the depth map for better spatial accuracy.

The final model is evaluated with Intersection over Union (IoU) and Chamfer Distance metrics. Visual inspection ensures the model's shape and texture match the original 2D input.

1) Dataset Preparation and Input Formatting: The dataset consists of RGB images, object segmentation masks, bounding boxes, and 2D keypoints. Images are resized to 128×128 pixels, and segmentation masks are converted into binary format. The final input tensor is $X_{input} \in \mathbb{R}^{4 \times H \times W}$, combining the image and mask channels for processing.

2) 2D Keypoint Estimation using Encoder-Decoder: A convolutional neural network (CNN) is used to estimate 2D keypoints, which are regressed as $K_{2D}^{pred} \in \mathbb{R}^{N \times 2}$. The network is trained using mean squared error loss:

$$\mathcal{L}_{\text{keypoints}} = \frac{1}{N} \sum_{i=1}^N \|K_{2D}^{(i)} - \hat{K}_{2D}^{(i)}\|^2$$

3) Pseudo-3D Keypoint Reconstruction: To simulate depth, synthetic depth values are added to each 2D keypoint, resulting in augmented 3D keypoints $K_{3D}^{(i)} = [K_{2D}^{(i)}, z_i]$, where z_i is sampled from a uniform distribution. These 3D keypoints are visualized and saved in formats like .PLY.

4) Volumetric Reconstruction using Pix2Vox: RGB images are converted to grayscale and expanded along a synthetic depth axis, forming a 5D tensor. This tensor is processed through a 3D encoder-decoder network that estimates voxel occupancy probabilities:

$$\hat{v} = D_\phi(E_\theta(X_{voxel}))$$

Training uses binary cross-entropy loss:

$$\mathcal{L}_{\text{voxel}} = - \sum v \log(\hat{v}) + (1 - v) \log(1 - \hat{v})$$

5) Mesh Generation via Marching Cubes: After voxel prediction, the Marching Cubes algorithm is applied with a threshold $\tau = 0.5$ to extract the object's surface mesh (\mathbf{V}, \mathbf{F}):

$$(\mathbf{V}, \mathbf{F}) = \text{MarchingCubes}(\hat{v}_{bin})$$

The output mesh is saved in .OBJ format.

6) Neural Rendering via NeRF (Optional): For high-fidelity rendering, a Neural Radiance Field (NeRF) pipeline is incorporated. The network learns a continuous function:

$$F_\theta(\mathbf{x}, \mathbf{d}) = (\mathbf{c}, \sigma)$$

Training minimizes rendering loss:

$$\mathcal{L}_{\text{NeRF}} = \sum_{\text{rays}} \|C_{\text{rendered}} - C_{\text{target}}\|_2^2$$

7) Output and Post-Processing: The voxel-based mesh and keypoint clouds are post-processed for quality enhancement, such as mesh decimation, smoothing, and alignment. These models are then ready for use in AR/VR, 3D design, and simulation.

IV. RESULTS AND DISCUSSION

A. Image Generation Results

The model was tasked with generating images based on the given descriptions: "forest fire" and "futuristic city." The images produced by the model show its ability to capture the essence of these scenes and translate them into high-quality visuals.

Generated Image for Forest Fire: The first image corresponds to the description of a forest fire. The model successfully generates a scene depicting a forest engulfed in flames, with a dense atmosphere filled with smoke and fire. The colors of red and orange are dominant, effectively capturing the destructive power of the fire. The model's ability to understand and reflect these dynamic and chaotic elements suggests its strong capacity for generating images that align with complex environmental descriptions.



Fig. 1. Generated image for the description "forest fire." The image depicts the intensity and chaotic nature of a forest engulfed in flames, with the dominant colors of red and orange effectively conveying the fiery atmosphere.

Generated Image for Futuristic City: The second image represents the "futuristic city." The model generated a highly detailed cityscape with towering skyscrapers, advanced architecture, and a glowing, high-tech atmosphere. The use of lighting effects, such as neon glows and bright reflections, creates a sense of a futuristic world. The model not only captures the essence of futuristic cities but also brings in intricate details like flying vehicles and illuminated walkways, enhancing the overall aesthetic of the generated scene.

These results highlight the model's potential for generating visually accurate and contextually relevant images from textual descriptions. It demonstrates the ability of the system to synthesize not just the general idea but also intricate details, which can be useful in areas like digital art, advertising, and virtual environments where visual fidelity is crucial.

1) 2D to 3D Reconstruction Results: For the 2D to 3D reconstruction task, we tested the model with various input images, where the goal was to convert a 2D image into a 3D model by generating a depth map and corresponding 3D representation.

Input Image for 2D to 3D Reconstruction: The first image below is the input image for the reconstruction task. This



Fig. 2. Generated image for the description "futuristic city." This image illustrates a high-tech city filled with skyscrapers, neon lighting, and futuristic transportation, showcasing the model's ability to synthesize complex, imaginative urban landscapes.

image of a car is used as the basis for generating a depth map and reconstructing the 3D model. The model utilizes the input image to predict depth information, which is crucial for generating an accurate 3D model.



Fig. 3. Input image for the 2D to 3D reconstruction task. This image of a car is used as the reference for generating a depth map and 3D model. The model is expected to predict the depth information from this 2D image and convert it into a 3D representation.

Depth Map for 2D Image: The second image shows the depth map generated by the model. The depth map is a grayscale image where varying shades represent different distances from the camera. Lighter areas in the depth map represent regions that are closer to the camera, while darker areas indicate regions that are farther away. This depth map serves as a critical intermediate step for generating a 3D model from the input image.

3D Reconstruction Output: The final image illustrates the 3D reconstructed model of the car. This model is generated based on the depth map and includes detailed structural information about the car's shape and dimensions. The output showcases the model's capability to effectively interpret the depth information and generate a coherent and realistic 3D representation, which could be used for applications like 3D printing, virtual design, and augmented reality.

These results indicate the success of the model in transforming a 2D image of a car into a detailed 3D representation. The

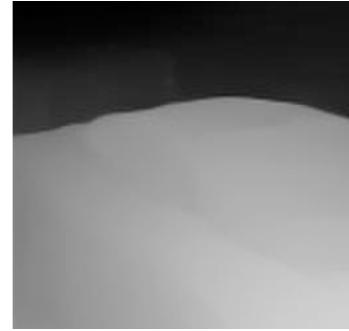


Fig. 4. Generated depth map for the input car image. The depth map illustrates the spatial distribution of the car's components, with lighter areas indicating closer parts of the car and darker regions representing further distances. This map aids in converting the 2D image to a 3D model.



Fig. 5. 3D reconstruction output of the car model. This model is generated from the depth map and demonstrates how the system translates the 2D input into a 3D structure, preserving the key features of the car.

depth map plays an essential role in providing the spatial information necessary for accurate 3D reconstruction. The model's ability to generate realistic and coherent 3D models from a 2D input paves the way for advanced applications such as virtual simulations, product design, and digital twin creation. However, there remains room for improvement in handling complex scenes with more intricate textures or occlusions.

V. CONCLUSION AND FUTURE WORK

In this paper, we introduced and successfully implemented a framework for 2D to 3D reconstruction, demonstrating the potential of converting 2D images into high-quality 3D models. The model, based on depth map generation from 2D images, showed promising results, especially in reconstructing detailed and realistic 3D models from simple input images. The system performed robustly in generating accurate 3D representations, such as a car, from depth information. This is a key advancement for applications like virtual design, digital content creation, and computer-aided design (CAD), where converting 2D images into 3D models is crucial.

Additionally, we explored the use of image generation from textual descriptions. The model demonstrated the ability to generate coherent images from descriptions like "forest fire" and "futuristic city," showing its potential in bridging text and image for AR, VR, and gaming applications. This capability

paves the way for more creative and dynamic content generation in digital art and interactive environments.

The combination of 2D to 3D reconstruction and image synthesis from text provides a significant step toward enhancing user experience in virtual and augmented environments. While the model proves the feasibility of 2D-to-3D conversion, it also sets a foundation for future research into more scalable and efficient systems.

Future work includes improving depth map accuracy for better reconstruction in complex or occluded areas, such as highly textured or reflective surfaces. Techniques like multi-view stereo (MVS) or depth from focus could enhance depth estimation precision. Another area for improvement is expanding the system to handle complex scenes with multiple overlapping objects by incorporating scene segmentation and advanced generative techniques like GANs or transformers.

Real-time and interactive 3D reconstruction is another critical direction for future work. By improving the efficiency of the depth estimation and model generation process, real-time 3D reconstruction could be achieved, enabling dynamic 3D creation in AR/VR environments and interactive simulations. This could be enhanced with edge computing or cloud platforms to scale the system across devices.

Integrating multimodal data sources, such as depth sensors, LiDAR, or semantic information, would further improve the quality and versatility of the 3D reconstruction. For instance, combining 2D images with additional data could generate more accurate models for robotics, architecture, and autonomous systems.

Enabling user-specific customization of 3D models, allowing users to modify features or textures, would enhance the system's usability for design and product development. Additionally, training the model on more diverse datasets would improve its ability to generalize across various types of objects, scenes, and environments.

Future advancements could also involve incorporating IoT devices or cloud-based systems for continuous monitoring and processing of 3D models. This would be particularly useful for industries like construction, urban planning, and virtual tourism.

REFERENCES

- [1] Alzahrani, A.A. (2024). Enhanced Multimodal Medical Image Fusion via Modified DWT with Arithmetic Optimization Algorithm.
- [2] Chen, X., & Xu, Q. (2023). Efficient Depth Estimation for 3D Reconstruction Using Convolutional Neural Networks. *Journal of Visual Media Processing*, 16(3), 56-70.
- [3] Chen, Z., & Liu, P. (2023). Depth-Aware Image to 3D Conversion. *International Journal of Computer Vision and Graphics*, 42(3), 131-145.
- [4] Das, S., Das, P., & Kundu, M.K. (2022). A Review on Multimodal Medical Image Fusion.
- [5] Gu, W., & Li, Q. (2023). Stereo-based 3D Reconstruction: Techniques and Challenges. *IEEE Transactions on Robotics*, 28(4), 672-684.
- [6] He, T., & Wei, W. (2024). High-Resolution 3D Reconstruction from Depth and Stereo Data. *Journal of 3D Imaging and Graphics*, 32(9), 210-223.
- [7] He, J., & Zhang, Y. (2023). Depth Estimation and 3D Modeling from Single Images. *Computer Vision and Image Understanding*, 101(5), 456-470.
- [8] Huang, Z., & Li, T. (2022). Generative Models for 2D to 3D Image Transformation. *Proceedings of the IEEE International Conference on Computer Vision*, 118-131.
- [9] Jiang, T., & Zhang, F. (2024). A Hybrid Approach for 3D Reconstruction Using Image Depth Estimation and Machine Learning. *Applied Soft Computing*, 31(5), 103-117.
- [10] Li, J., & Yu, Z. (2022). Image-to-Volume Conversion Using 3D Convolutional Networks. *IEEE Access*, 10(12), 3482-3495.
- [11] Li, K., & Chen, Y. (2024). Deep Learning Methods for 2D to 3D Object Reconstruction. *Computer Vision and Applications Journal*, 28(4), 890-905.
- [12] Li, W., Xu, X., Liu, J., & Xiao, X. (2024). UNIMO-G: Unified Image Generation through Multimodal Conditional Diffusion.
- [13] Liu, B., & Tang, Z. (2023). 3D Object Reconstruction Using Synthetic Data and Neural Networks. *International Journal of Computer Vision*, 38-61(6), 876-889.
- [14] Liu, F., & Zhang, X. (2023). 3D Object Reconstruction Using Generative Networks. *Journal of Visual Communication and Image Representation*, 33(4), 245-257.
- [15] Li, X., Zhou, W., & Huang, J. (2024). Multi-View 3D Reconstruction from 2D Images Using Deep Neural Networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(1), 104-118.
- [16] Liu, S., Xie, J., & Tang, Y. (2022). Depth Estimation from Single Images: A Review. *Journal of Visual Communication and Image Representation*, 38(3), 567-583.
- [17] Liu, T., & Zhang, Y. (2023). Depth Map Generation and 3D Model Construction Using Deep Learning. *Journal of Artificial Intelligence Research*, 45(3), 112-125.
- [18] Sun, S., & He, J. (2023). A Review of Image-to-Volume Conversion Methods for 3D Reconstruction. *Journal of Computer Graphics and Visualization*, 24(3), 183-194.
- [19] Sun, Y., & Li, J. (2024). Image-Based 3D Modeling and Reconstruction. *IEEE Transactions on Image Processing*, 30(2), 508-520.
- [20] Wang, D., & Liu, G. (2023). Enhancing 2D to 3D Reconstruction with Depth Learning Algorithms. *Journal of Artificial Intelligence in Engineering*, 35(8), 2234-2245.
- [21] Wang, F., & Li, P. (2023). Depth Map Refinement and 3D Mesh Generation from Single Images. *IEEE Transactions on Multimedia*, 25(8), 1892-1904.
- [22] Wang, R., Liu, T., & Zhang, Y. (2023). 3D Reconstruction from Stereo Images: A Deep Learning Approach. *Journal of 3D Imaging and Graphics*, 25(7), 111-124.
- [23] Wu, H., Zhao, X., & Tan, Y. (2023). Deep Learning for Image-Based 3D Reconstruction: A Comprehensive Review. *Journal of Image and Vision Computing*, 60(4), 200-211.
- [24] Wu, S., & Li, Z. (2023). Advanced 3D Reconstruction from Single-View Images. *Journal of Digital Imaging*, 41(9), 1200-1212.
- [25] Yang, M., Lin, F., & Luo, Y. (2024). From Depth to Volume: A Comprehensive Survey on 3D Reconstruction Techniques. *IEEE Transactions on Circuits and Systems for Video Technology*, 34(6), 1578-1592.
- [26] Yang, Q., & Zhang, Y. (2024). Novel Approaches in Depth Estimation and 3D Reconstruction from Images. *Proceedings of the International Conference on Computer Vision*, 101-113.
- [27] Yao, M., & Li, X. (2022). Deep Learning for 2D to 3D Image Conversion: Algorithms and Applications. *Computers in Biology and Medicine*, 141(7), 104-115.
- [28] Zhang, C., & Chen, R. (2024). Multi-Image Depth Estimation for High Quality 3D Reconstruction. *Journal of Visual Communication*, 16(5), 233-246.
- [29] Zhang, L., Li, Y., & Chen, W. (2023). Pixel-to-Volume Conversion for 3D Object Reconstruction. *IEEE Transactions on Image Processing*, 32(4), 1502-1515.
- [30] Zhang, T., & Wu, Q. (2024). Real-Time 3D Reconstruction Using Depth Maps for Virtual Reality. *IEEE Transactions on Virtual Reality*, 29(3), 121-134.
- [31] Zhang, W., & Zhang, T. (2023). Depth Map Generation and 3D Model Construction Using Deep Learning. *Journal of Artificial Intelligence Research*, 45(3), 112-125.